

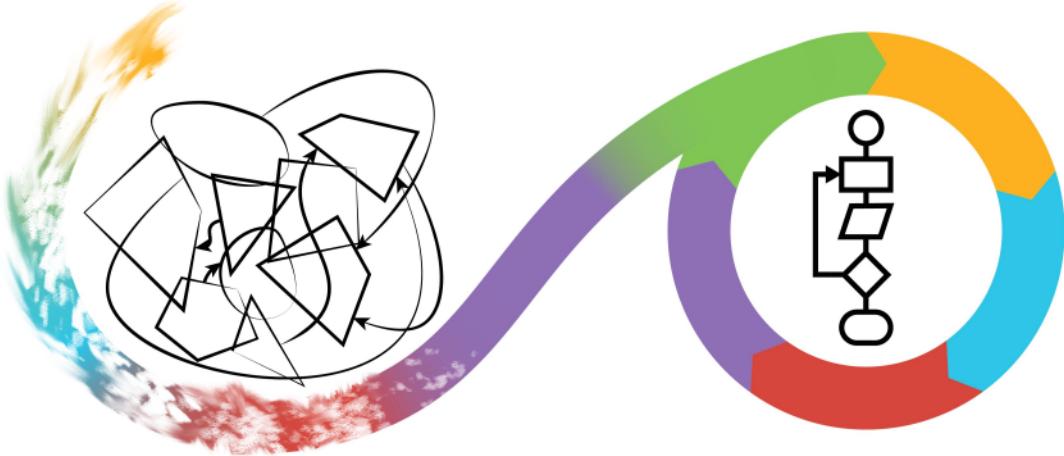
# Optimisation

Richèl Bilderbeek

## 1 The Big Picture



[https://github.com/UPPMAX/programming\\_formalisms/blob/main/tdd/tdd\\_lecture/tdd\\_lecture.qmd](https://github.com/UPPMAX/programming_formalisms/blob/main/tdd/tdd_lecture/tdd_lecture.qmd)



### 1.1 Breaks

Please take breaks: these are important for learning. Ideally, do something boring (1)!

### 1.2 Schedule

Day	From	To	What
Fri	12:00	13:00	Lunch
Fri	13:00	14:00	Lecture: optimization
Fri	14:00	14:15	Break

Day	From	To	What
Fri	14:15	15:00	Project: write/optimize code
Fri	15:00	15:15	Break
Fri	15:15	15:45	Project: write/optimize code
Fri	15:45	16:00	Reflection

## 2 Optimisation

[https://github.com/UPPMAX/programming\\_formalisms/blob/main/optimisation/optimisation\\_lecture/optimisation\\_lecture.qmd](https://github.com/UPPMAX/programming_formalisms/blob/main/optimisation/optimisation_lecture/optimisation_lecture.qmd)



## 3 Why optimization?

To improve the runtime speed (or memory use) of a program



Captain Obvious

## 4 Misconceptions

Q: What would be **bad advice** to improve the run-time speed of an algorithm?

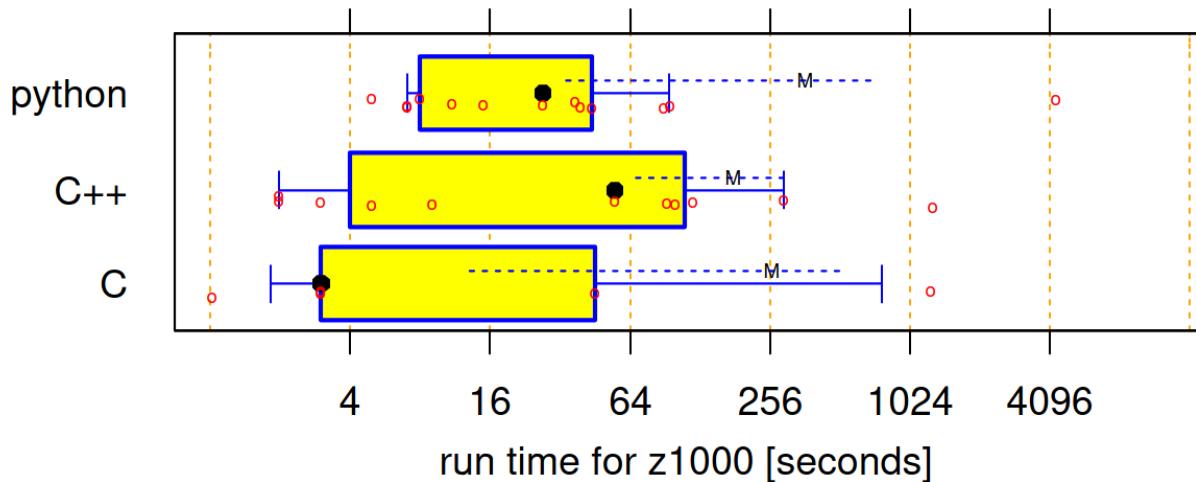
Fill in in the shared document!

### 4.1 Bad advice 1

'Use C or C++ or Rust'

...

Variance within programming languages is bigger than variance between languages (adapted fig 2, from (2))



#### 4.2 Bad advice 2

'No for loops', 'unroll for-loops', any other micro-optimization.

...

Premature optimization is the root of all evil. It likely has no measurable effect.

#### 4.3 Bad advice 2

We should forget about small efficiencies, say about 97% of the time: premature optimization is the root of all evil. Yet we should not pass up our opportunities in that critical 3%.

Donald Knuth

Source: [Wikipedia](#)

#### 4.4 Bad advice 3

'Always parallelize'

...

- Maximum gain depends on proportion spent in the parallelized part (3)
- Overhead is underestimated
- Hard to debug



Figure 1: Donald Knuth

#### 4.5 Bad advice 3

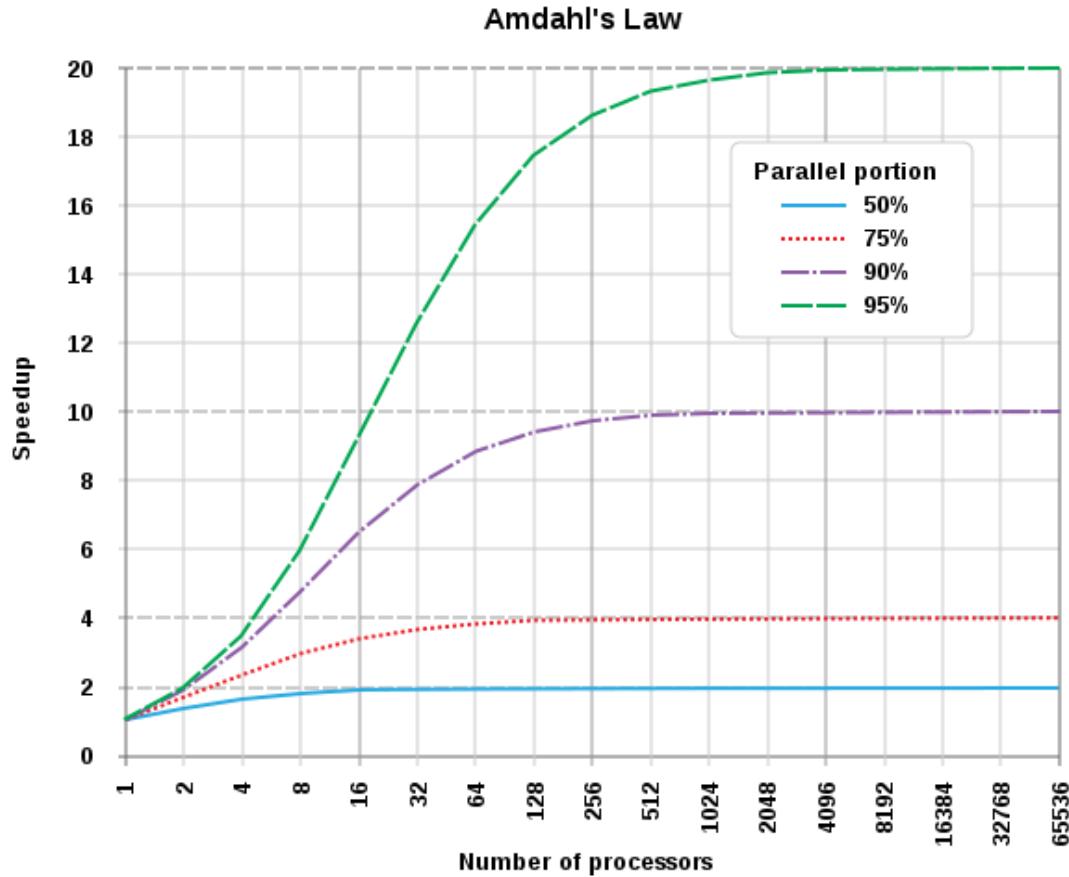


Figure 2: <https://en.wikipedia.org/wiki/File:AmdahlsLaw.svg#file>

#### 4.6 Bad advice 4

'Optimize the function where you feel the performance problem is'

Developers -also very experienced developers- are known to have a bad intuition (4)

Instead, from (5):

1. finding code program spends most time in
2. measure timing of that code
3. analyze the measured runtimes

## 4.7 Bad advice 5

'Optimize each function'

- The 90-10 rule: 90% of all time, the program spends in 10% of the code.
- Your working hours can be spent once

## 5 Proper method

### 5.1 Problem

Q: When to optimize for speed?

...

A:

- C++ Core Guidelines: Per.1: Don't optimize without reason
- C++ Core Guidelines: Per.2: Don't optimize prematurely
- C++ Core Guidelines: Per.3: Don't optimize something that's not performance critical

### 5.2 Problem

Q: How to improve the run-time speed of an algorithm?

...

Make it work, make it right, make it fast.

Kent Beck

A (simplified):

1. Measure (hard to do (6))
2. Think
3. Change code
4. Measure again

### **5.3 Problem**

Q: How to improve the run-time speed of an algorithm?

A (simplified):

1. Measure big-O
2. Measure speed profile
3. Think
4. Change code
5. Measure again

### **5.4 Measurement 1: big-O**

How your (combination of) algorithms scales with more complex input.

- Counting the words in a book:  $O(n)$
- Looking up a word in a dictionary:  $O(\log_2(n))$

Do measure big-O in release mode!

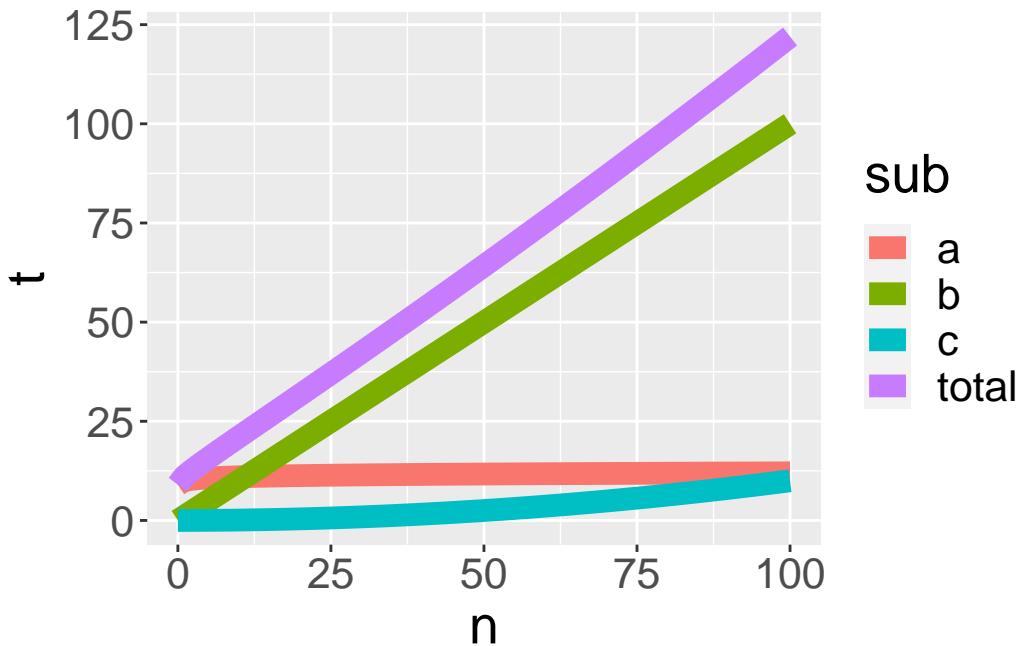
## 5.5 Your algorithm



## 5.6 Example

```
create_big_o_example <- function(n = seq(0, 100)) {  
  t_wide <- tibble::tibble(n = n)  
  t_wide$a <- 10 + log10(t_wide$n + 0.1)  
  t_wide$b <- t_wide$n  
  t_wide$c <- 0.001 * (t_wide$n ^ 2)  
  t_wide$total <- t_wide$a + t_wide$b + t_wide$c  
  t <- tidyr::pivot_longer(t_wide, cols = c("a", "b", "c", "total"))  
  colnames(t) <- c("n", "sub", "t")  
  t  
}  
t <- create_big_o_example(n = seq(0, 100))  
ggplot2::ggplot(t, ggplot2::aes(x = n, y = t, color = sub)) +  
  ggplot2::geom_line(size = 4) +  
  ggplot2::theme(text = ggplot2::element_text(size = 20))
```

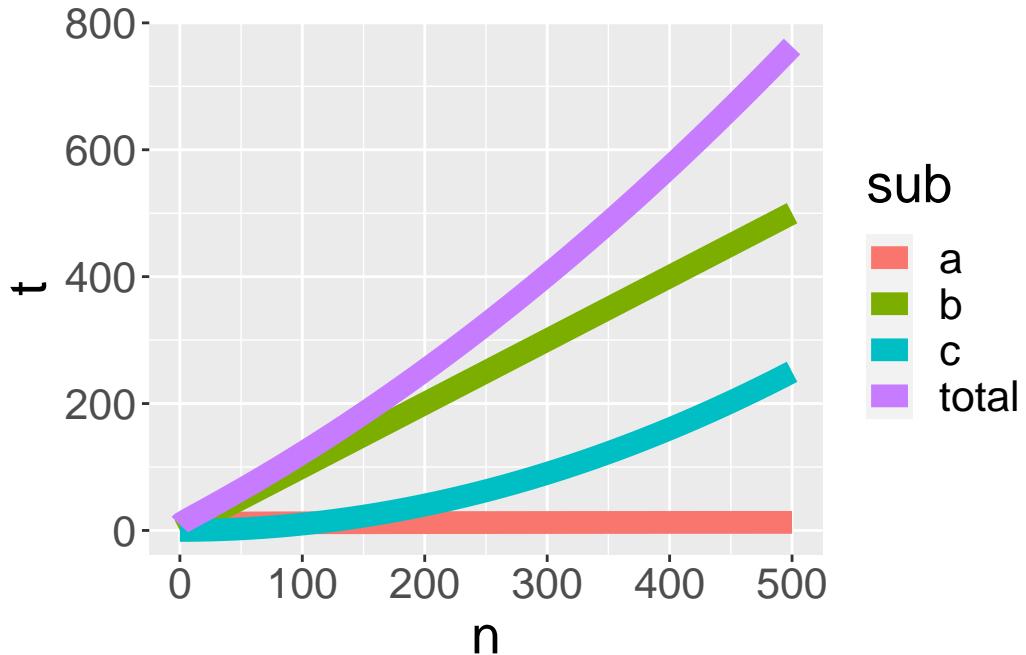
Warning: Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.  
i Please use `linewidth` instead.



Work on B?

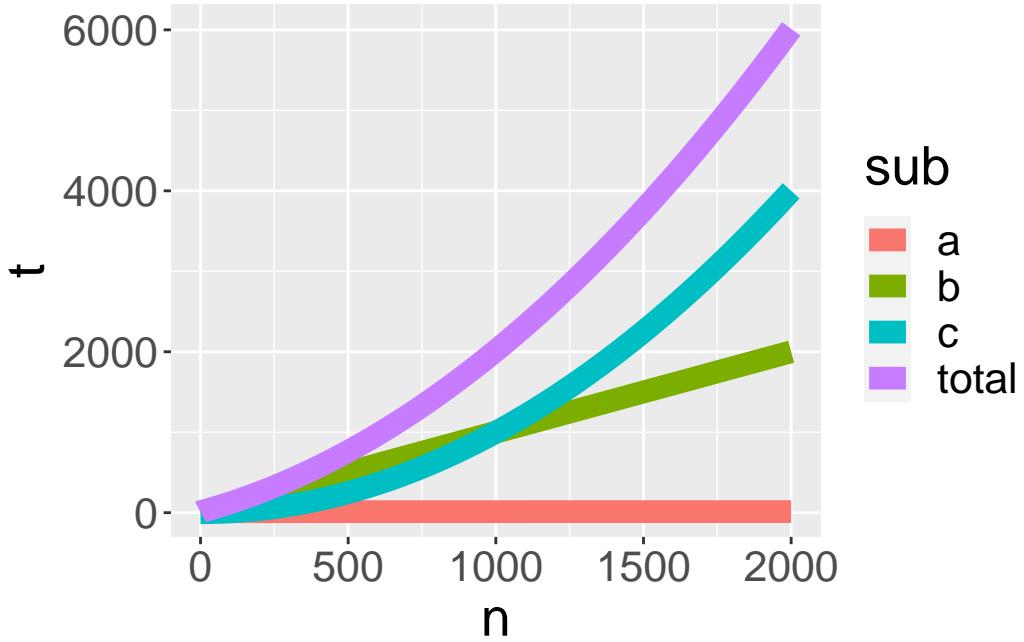
## 5.7 Example

```
t <- create_big_o_example(n = seq(0, 500))
ggplot2::ggplot(t, ggplot2::aes(x = n, y = t, color = sub)) +
  ggplot2::geom_line(size = 4) +
  ggplot2::theme(text = ggplot2::element_text(size = 20))
```



## 5.8 Example

```
t <- create_big_o_example(n = seq(0, 2000))
ggplot2::ggplot(t, ggplot2::aes(x = n, y = t, color = sub)) +
  ggplot2::geom_line(size = 4) +
  ggplot2::theme(text = ggplot2::element_text(size = 20))
```



No, work on C instead

## 5.9 Conclusions

Big-O helps to:

- find algorithm to profile
- make predictions

## 5.10 Exercise 1

- Measure big-O complexity of <https://www.pythontutorial.net/python-basics/python-prime-number/>

```
def isprime(num):
    for n in range(
        2, int(num**0.5)+1
    ):
        if num%n==0:
```

```

        return False
    return True

def isprime(num):
    if num > 1:
        for n in range(2,num):
            if (num % n) == 0:
                return False
        return True
    else:
        return False

```

### 5.11 Exercise 1

- Measure big-O complexity of <https://www.pythontutorial.net/python-basics/python-prime-number/>

```

def isprime(num):
    for n in range(
        2, int(num**0.5)+1
    ):
        if num%n==0:
            return False
    return True

```

```

def Prime(no, i = 2):
    if no == i:
        return True
    elif no % i == 0:
        return False
    return Prime(no, i + 1)

```

## 5.12 Exercise 2

- Measure big-O complexity of DNA alignment at <https://johnlekberg.com/blog/2020-10-25-seq-align.html>

```
ACGTACGTACGTACGTACGTACGT  
ACGTACGTACGTCGTACGTACGT
```

```
ACGTACGTACGTACGTACGTACGT  
ACGTACGTACGT-CGTACGTACGT
```

## 5.13 Measurement 2: Run-time speed profile

- See which code is spent most time in
  - Use an input of suitable complexity
    - Note to self: next example should take at least 10 seconds!
- Consider using CI to obtain a speed profile every push!

## 5.14 Run-time speed profile: code

```
def is_sorted(data): return data == sorted(data)
def richel_sort(data):
    while not is_sorted(data):
        from random import shuffle
        shuffle(data)
    return data

def create(): return list(range(0, 10))
def reverse(x):
    x.reverse()
    return x
def sort(x): return richel_sort(x)

def my_function(): return sort(reverse(create()))

import cProfile
cProfile.run("my_function()")
```

## 5.15 Run-time speed profile: results

149384793 function calls (149384751 primitive calls) in 44.759 seconds

Ordered by: standard name

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
6	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap>:100(acquire)
6/1	0.000	0.000	0.002	0.002	<frozen importlib._bootstrap>:1022(_find_and_load)
6	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap>:125(release)
6	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap>:165(__init__)
6	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap>:169(__enter__)
6	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap>:173(__exit__)
6	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap>:179(_get_module_by_name)
6	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap>:198(cb)
10/1	0.000	0.000	0.002	0.002	<frozen importlib._bootstrap>:233(_call_with_frameskip)
140	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap>:244(_verbose_message)
6	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap>:357(__init__)
8	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap>:391(cached)
6	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap>:404(parent)
6	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap>:412(has_location)
2	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap>:48(__new_module)
6	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap>:492(__init_module)
6	0.000	0.000	0.001	0.000	<frozen importlib._bootstrap>:564(module_from_spec)
6/1	0.000	0.000	0.002	0.002	<frozen importlib._bootstrap>:664(_load_unlocked)
6	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap>:71(__init__)
6	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap>:746(find_spec)
6	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap>:826(find_spec)
24	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap>:893(__enter__)
24	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap>:897(__exit__)
6	0.000	0.000	0.001	0.000	<frozen importlib._bootstrap>:921(_find_spec)
6/1	0.000	0.000	0.002	0.002	<frozen importlib._bootstrap>:987(_find_and_load)
2	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap_external>:1040(__init__)
2	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap_external>:1065(get_code)
2	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap_external>:1070(get_source)
2	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap_external>:1089(open)
4	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap_external>:1163(__enter__)
4	0.000	0.000	0.001	0.000	<frozen importlib._bootstrap_external>:1174(create)
4	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap_external>:1182(execute)
126	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap_external>:126(__path__)
126	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap_external>:128(<listcomp>, 0)
4	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap_external>:132(__path__)

8	0.000	0.000	0.000	0.000 <frozen importlib._bootstrap_external>:134(<gen
34	0.000	0.000	0.000	0.000 <frozen importlib._bootstrap_external>:1356(_pa
6	0.000	0.000	0.001	0.000 <frozen importlib._bootstrap_external>:1399(_g
36	0.000	0.000	0.000	0.000 <frozen importlib._bootstrap_external>:140(_pa
6	0.000	0.000	0.001	0.000 <frozen importlib._bootstrap_external>:1431(fir
6	0.000	0.000	0.000	0.000 <frozen importlib._bootstrap_external>:150(_pa
6	0.000	0.000	0.000	0.000 <frozen importlib._bootstrap_external>:1531(_g
28	0.000	0.000	0.001	0.000 <frozen importlib._bootstrap_external>:1536(fir
6	0.000	0.000	0.000	0.000 <frozen importlib._bootstrap_external>:159(_pa
6	0.000	0.000	0.000	0.000 <frozen importlib._bootstrap_external>:180(_pa
4	0.000	0.000	0.000	0.000 <frozen importlib._bootstrap_external>:380(cach
6	0.000	0.000	0.000	0.000 <frozen importlib._bootstrap_external>:510(_get
2	0.000	0.000	0.000	0.000 <frozen importlib._bootstrap_external>:542(_ch
2	0.000	0.000	0.000	0.000 <frozen importlib._bootstrap_external>:585(_cl
2	0.000	0.000	0.000	0.000 <frozen importlib._bootstrap_external>:618(_va
2	0.000	0.000	0.000	0.000 <frozen importlib._bootstrap_external>:670(_co
28	0.000	0.000	0.000	0.000 <frozen importlib._bootstrap_external>:71(_rela
6	0.000	0.000	0.000	0.000 <frozen importlib._bootstrap_external>:721(spec
6	0.000	0.000	0.000	0.000 <frozen importlib._bootstrap_external>:84(unpa
2	0.000	0.000	0.000	0.000 <frozen importlib._bootstrap_external>:874(crea
2/1	0.000	0.000	0.002	0.002 <frozen importlib._bootstrap_external>:877(exe
2	0.000	0.000	0.000	0.000 <frozen importlib._bootstrap_external>:950(get
1	0.000	0.000	44.759	44.759 <string>:1(<module>)
4060761	0.975	0.000	2.754	0.000 <string>:1(is_sorted)
	1	0.000	0.000	0.000 <string>:1(reverse)
	1	0.000	0.000	44.759 <string>:1(sort)
	1	4.605	4.605	44.759 <string>:1(stupid_sort)
	1	0.000	0.000	0.000 <string>:2(create)
	1	0.000	0.000	44.759 <string>:2(my_function)
	6	0.000	0.000	0.000 __init__.py:89(find_spec)
	6	0.000	0.000	0.000 __init__.py:96(<lambda>)
	1	0.000	0.000	0.000 bisect.py:1(<module>)
	6	0.000	0.000	0.000 call.py:5(python_function)
6/1	0.000	0.000	0.003	0.003 loader.py:113(_find_and_load_hook)
6/1	0.000	0.000	0.002	0.002 loader.py:115(_hook)
	6	0.000	0.000	0.000 loader.py:40(_maybe_run_hooks)
6/1	0.000	0.000	0.003	0.003 loader.py:82(_run_hook)
	1	0.000	0.000	0.002 random.py:1(<module>)
	1	0.000	0.000	0.000 random.py:103(Random)
	1	0.000	0.000	0.000 random.py:119(__init__)
	1	0.000	0.000	0.000 random.py:128(seed)
	1	0.000	0.000	0.000 random.py:219(__init_subclass__)
36546840	18.451	0.000	23.971	0.000 random.py:239(_randbelow_with_getrandbits)

4060760	13.122	0.000	37.397	0.000 random.py:380(shuffle)
1	0.000	0.000	0.000	0.000 random.py:813(SystemRandom)
1	0.000	0.000	0.000	0.000 threading.py:1430(current_thread)
2	0.000	0.000	0.000	0.000 {built-in method _imp._fix_co_filename}
36	0.000	0.000	0.000	0.000 {built-in method _imp.acquire_lock}
4	0.001	0.000	0.001	0.000 {built-in method _imp.create_dynamic}
4	0.000	0.000	0.000	0.000 {built-in method _imp.exec_dynamic}
6	0.000	0.000	0.000	0.000 {built-in method _imp.is_builtin}
6	0.000	0.000	0.000	0.000 {built-in method _imp.is_frozen}
36	0.000	0.000	0.000	0.000 {built-in method _imp.release_lock}
12	0.000	0.000	0.000	0.000 {built-in method _thread.allocate_lock}
13	0.000	0.000	0.000	0.000 {built-in method _thread.get_ident}
2	0.000	0.000	0.000	0.000 {built-in method builtins.__build_class__}
3/1	0.000	0.000	44.759	44.759 {built-in method builtins.exec}
42	0.000	0.000	0.000	0.000 {built-in method builtinsgetattr}
41	0.000	0.000	0.000	0.000 {built-in method builtins.hasattr}
41	0.000	0.000	0.000	0.000 {built-in method builtins.isinstance}
4060768	0.304	0.000	0.304	0.000 {built-in method builtins.len}
6	0.000	0.000	0.000	0.000 {built-in method builtins.locals}
4	0.000	0.000	0.000	0.000 {built-in method builtins.max}
4060761	1.779	0.000	1.779	0.000 {built-in method builtins.sorted}
6	0.000	0.000	0.000	0.000 {built-in method from_bytes}
2	0.000	0.000	0.000	0.000 {built-in method io.open_code}
2	0.000	0.000	0.000	0.000 {built-in method marshal.loads}
1	0.000	0.000	0.000	0.000 {built-in method math.exp}
2	0.000	0.000	0.000	0.000 {built-in method math.log}
1	0.000	0.000	0.000	0.000 {built-in method math.sqrt}
10	0.000	0.000	0.000	0.000 {built-in method posix.fspath}
6	0.000	0.000	0.000	0.000 {built-in method posix.getcwd}
1	0.000	0.000	0.000	0.000 {built-in method posix.register_at_fork}
36	0.000	0.000	0.000	0.000 {built-in method posix.stat}
6	0.000	0.000	0.000	0.000 {built-in method rpycall.call_r_function}
1	0.000	0.000	0.000	0.000 {function Random.seed at 0x7ffb9a2870a0}
2	0.000	0.000	0.000	0.000 {method '__exit__' of '_io._IOBase' objects}
12	0.000	0.000	0.000	0.000 {method '__exit__' of '_thread.lock' objects}
6	0.000	0.000	0.000	0.000 {method 'append' of 'list' objects}
36546840	2.027	0.000	2.027	0.000 {method 'bit_length' of 'int' objects}
1	0.000	0.000	0.000	0.000 {method 'disable' of '_lsprof.Profiler' objects}
10	0.000	0.000	0.000	0.000 {method 'endswith' of 'str' objects}
6	0.000	0.000	0.000	0.000 {method 'format' of 'str' objects}
12	0.000	0.000	0.000	0.000 {method 'get' of 'dict' objects}
60046373	3.493	0.000	3.493	0.000 {method 'getrandbits' of '_random.Random' objects}
130	0.000	0.000	0.000	0.000 {method 'join' of 'str' objects}

6	0.000	0.000	0.000	0.000 {method 'pop' of 'dict' objects}
2	0.000	0.000	0.000	0.000 {method 'read' of '_io.BufferedReader' objects}
1	0.000	0.000	0.000	0.000 {method 'reverse' of 'list' objects}
4	0.000	0.000	0.000	0.000 {method 'rfind' of 'str' objects}
44	0.000	0.000	0.000	0.000 {method 'rpartition' of 'str' objects}
256	0.000	0.000	0.000	0.000 {method 'rstrip' of 'str' objects}
6	0.000	0.000	0.000	0.000 {method 'startswith' of 'str' objects}

## 5.16 Myth 1

```

def slow_tmp_swap(x, y):
    tmp = x
    x = y
    y = tmp
    return x, y

def superfast_xor_swap(x, y):
    x ^= y
    y ^= x
    x ^= y
    return x, y

```

..

- C++ Core Guidelines: Per.4: Don't assume that complicated code is necessarily faster than simple code
- C++ Core Guidelines: Per.5: Don't assume that low-level code is necessarily faster than high-level code

## 5.17 Exercise

Create speed profile of <https://www.pythontutorial.net/python-basics/python-speed-profile/>

Create speed profile of DNA alignment

## 5.18 Step 3: Think

- How to achieve the same with less calculations?
  - Aim to change big-O, not some micro-optimization
  - For example, store earlier results in a sorted look-up table

Feynman Problem Solving Algorithm:

1. Write down the problem.
2. Think very hard.
3. Write down the answer

### 5.19 Step 4: Measure again

In TDD, this test would have been present already:

```
assert 10.0 * get_t_runtime_b() < get_t_runtime_a()
```

Adapt the constant to reality.

- C++ Core Guidelines: Per.6: Don't make claims about performance without measurements

### 5.20 Recap quote

It is far, far easier to make a correct program fast, than it is to make a fast program correct.

Herb Sutter

Source [Wikimedia](#)

### 5.21 Recap

- Be critical on speed optimization solutions
- Tested and clean code always comes first
- Measure correctly, at the right complexity, before and after
- Prefer changing big-O over micro-optimizations (but see first point!)

### 5.22 The End

### 5.23 Links

- Lecture of 2022: [here](#):



Figure 3: Herb Sutter

1. Newport C. Deep work: Rules for focused success in a distracted world. Hachette UK; 2016.
2. Prechelt L. An empirical comparison of c, c++, java, perl, python, rexx and tcl. *IEEE Computer*. 2000;33(10):23–9.
3. Rodgers DP. Improvements in multiprocessor system design. *ACM SIGARCH Computer Architecture News*. 1985;13(3):225–31.
4. Sutter H, Alexandrescu A. C++ coding standards: 101 rules, guidelines, and best practices. Pearson Education; 2004.
5. Chellappa S, Franchetti F, Püschel M. How to write fast numerical code: A small introduction. Generative and Transformational Techniques in Software Engineering II: International Summer School, GTTSE 2007, Braga, Portugal, July 2-7, 2007 Revised Papers. 2008;196–259.
6. Bartz-Beielstein T, Doerr C, Berg D van den, Bossek J, Chandrasekaran S, Eftimov T, et al. Benchmarking in optimization: Best practice and open issues. arXiv preprint arXiv:200703488. 2020;